

Cmpt 412

Project 3

Hao Ran Wei

301354701

Prereqs:

Turn on the following setting on google collab

### Notebook settings

#### Hardware accelerator

GPU  

#### GPU class

Standard 

#### Runtime shape

High-RAM 

Omit code cell output when saving this notebook

[Cancel](#) [Save](#)

---

## Part1:

List of final configs and modifications used:

Max iteration: 800  
Base LR: 0.0007  
Batch size: 512  
COCO model: faster\_rcnn\_X\_101\_32x8d\_FPN\_3x.yaml

Factors to help improve performance:

The reason why I chose the model faster\_rcnn\_X\_101\_32x8d\_FPN\_3x.yaml cause on the coco documentation the box AP had a higher accuracy than the default as shown in the ablation study. Also played mostly with base LR and max iteration to improve AP, due to both having the most impact on a total loss. Which is a good indicator of over fitting. I also opted to use the default data loader which had ResizeShortestEdge(short\_edge\_length=(800, 800), max\_size=1333, sample\_style='choice') as its augmentation.

### ablation study

Default

Config:

```
cfg.DATA_LOADER.NUM_WORKERS = 2
cfg.MODEL.WEIGHTS =
model_zoo.get_checkpoint_url("COCO-Detection/faster_rcnn_R_101_FPN_3x.yaml") # Let training initialize from model zoo
cfg.SOLVER.IMS_PER_BATCH = 2 # This is the real "batch size" commonly known to deep learning people
cfg.SOLVER.BASE_LR = 0.00025 # pick a good LR
cfg.SOLVER.MAX_ITER = 500 # 300 iterations seems good enough for this toy dataset; you will need to train Longer for a practical dataset
cfg.SOLVER.STEPS = [] # do not decay learning rate
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 512 # The "RoIHead batch size". 128 is faster, and good enough for this toy dataset (default: 512)
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 1
```

results:

```

creating index...
index created!
[02/25 19:38:59 d2.evaluation.fast_eval_api]: Evaluate annotation type "bbox"
[02/25 19:39:00 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.10 seconds.
[02/25 19:39:00 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[02/25 19:39:00 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.01 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.190
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.317
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.206
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.103
Average Precision (AP) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.252
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.500
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.013
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.106
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.213
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.101
Average Recall (AR) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.275
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.630
[02/25 19:39:00 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APM | API |
|:----:|:----:|:----:|:----:|:----:|:----:|
| 19.031 | 31.710 | 20.650 | 10.302 | 25.232 | 49.989 |
OrderedDict([('bbox', {'AP': 19.030540659393396, 'AP50': 31.709733096818766, 'AP75': 20.649587962667823, 'APs': 10.30200518759837, 'APM': 25.23187301024759, 'API': 49.988005599044994}))
```

List of the configs and modifications that you used.

Name	lr_sched	train_time (s/iter)	inference_time (s/im)	train_mem (GB)	box_AP	model_id	download
------	----------	---------------------	-----------------------	----------------	--------	----------	----------

Changed the default: [faster\\_rcnn\\_R\\_101\\_FPN\\_3x.yaml](#)

R101-FPN	3x	0.286	0.051	4.1	42.0	137851257	<a href="#">model</a>   <a href="#">metrics</a>
----------	----	-------	-------	-----	------	-----------	---

To: [faster\\_rcnn\\_X\\_101\\_32x8d\\_FPN\\_3x.yaml](#)

X101-FPN	3x	0.638	0.098	6.7	43.0	139173657	<a href="#">model</a>   <a href="#">metrics</a>
----------	----	-------	-------	-----	------	-----------	---

Reason: better box AP ratio

Results:

```

creating index...
index created!
[02/25 20:17:44 d2.evaluation.fast_eval_api]: Evaluate annotation type "bbox"
[02/25 20:17:44 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.13 seconds.
[02/25 20:17:44 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[02/25 20:17:44 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.01 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.240
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.393
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.272
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.138
Average Precision (AP) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.317
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.565
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.014
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.117
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.264
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.136
Average Recall (AR) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.344
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.669
[02/25 20:17:44 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APM | API |
|:----:|:----:|:----:|:----:|:----:|:----:|
| 23.953 | 39.337 | 27.188 | 13.770 | 31.740 | 56.542 |
OrderedDict([('bbox', {'AP': 23.95265285759814, 'AP50': 39.33731045946842, 'AP75': 27.187713550379446, 'APs': 13.770489019991519, 'APM': 31.740071153988225, 'API': 56.541637913252984}))
```

A much better results with an AP of 23.95 compared to 19.03.

Increase BASE\_LR to 0.0007

Results:

```

creating index...
index created!
[02/25 21:05:46 d2.evaluation.fast_eval_api]: Evaluate annotation type "bbox"
[02/25 21:05:46 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.16 seconds.
[02/25 21:05:46 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[02/25 21:05:46 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.01 seconds.
    Average Precision (AP) @| IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.314
    Average Precision (AP) @| IoU=0.50 | area= all | maxDets=100 ] = 0.310
    Average Precision (AP) @| IoU=0.75 | area= all | maxDets=100 ] = 0.351
    Average Precision (AP) @| IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.204
    Average Precision (AP) @| IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.403
    Average Precision (AP) @| IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.617
    Average Recall (AR) @| IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.016
    Average Recall (AR) @| IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.129
    Average Recall (AR) @| IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.342
    Average Recall (AR) @| IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.205
    Average Recall (AR) @| IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.435
    Average Recall (AR) @| IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.722
[02/25 21:05:46 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | AP1 |
|-----|-----|-----|-----|-----|-----|
| 31.425 | 51.022 | 35.323 | 20.351 | 40.336 | 61.746 |
OrderedDict([('bbox', {'AP': 31.424904330687873, 'AP50': 51.021593938926245, 'AP75': 35.323079223315396, 'APs': 20.351324180987163, 'APm': 40.33559901856648, 'AP1': 61.746205269859104}))]

```

**Significant increase from AP 23.95 to 31.**

Change batch size from 512 to 128 as suggested in the tutorial to speed up training speed.

**Result:**

```

~~~~ \~~~\~~~\~~~
creating index...
index created!
[02/25 21:39:04 d2.evaluation.fast_eval_api]: Evaluate annotation type "bbox"
[02/25 21:39:04 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.13 seconds.
[02/25 21:39:04 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[02/25 21:39:04 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.01 seconds.
    Average Precision (AP) @| IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.292
    Average Precision (AP) @| IoU=0.50 | area= all | maxDets=100 ] = 0.456
    Average Precision (AP) @| IoU=0.75 | area= all | maxDets=100 ] = 0.346
    Average Precision (AP) @| IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.178
    Average Precision (AP) @| IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.383
    Average Precision (AP) @| IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.617
    Average Recall (AR) @| IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.016
    Average Recall (AR) @| IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.126
    Average Recall (AR) @| IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.317
    Average Recall (AR) @| IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.178
    Average Recall (AR) @| IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.410
    Average Recall (AR) @| IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.720
[02/25 21:39:04 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | AP1 |
|-----|-----|-----|-----|-----|-----|
| 29.244 | 45.565 | 34.620 | 17.757 | 38.267 | 61.746 |
OrderedDict([('bbox', {'AP': 29.2435995514254, 'AP50': 45.56526556606061, 'AP75': 34.61980749098994, 'APs': 17.757068940214214, 'APm': 38.267113536313616, 'AP1': 61.746205269859104}))]

```

Notice how a lower batch size actually decreased AP, but barely.

Changing batch size from 512 to 768 to try to test AP.

**Results:**

```

DON (t=0.015)
creating index...
index created!
[02/25 22:28:27 d2.evaluation.fast_eval_api]: Evaluate annotation type "bbox"
[02/25 22:28:27 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.13 seconds.
[02/25 22:28:27 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[02/25 22:28:27 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.01 seconds.
    Average Precision (AP) @| IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.301
    Average Precision (AP) @| IoU=0.50 | area= all | maxDets=100 ] = 0.483
    Average Precision (AP) @| IoU=0.75 | area= all | maxDets=100 ] = 0.345
    Average Precision (AP) @| IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.195
    Average Precision (AP) @| IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.383
    Average Precision (AP) @| IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.602
    Average Recall (AR) @| IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.016
    Average Recall (AR) @| IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.128
    Average Recall (AR) @| IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.330
    Average Recall (AR) @| IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.195
    Average Recall (AR) @| IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.419
    Average Recall (AR) @| IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.721
[02/25 22:28:27 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | AP1 |
|-----|-----|-----|-----|-----|-----|
| 30.085 | 48.270 | 34.523 | 19.471 | 38.331 | 60.163 |
OrderedDict([('bbox', {'AP': 30.085400610910295, 'AP50': 48.28964491340521, 'AP75': 34.523160825362396, 'APs': 19.47108024072006, 'APm': 38.330874125800044, 'AP1': 60.16264288815803}))]

```

This seems not to affect the “AP” that much so I kept it as 512

Next increase the epoch from 500 to 800 since that's where the total loss seems to stagnate:

**Results:**

```
[02/27 20:24:15 d2.evaluation.coco_evaluation]: Evaluation results for bbox:  
| AP | AP50 | AP75 | APs | APM | AP1 |  
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:  
| 31.278 | 50.289 | 35.591 | 19.593 | 49.661 | 66.812 |  
OrderedDict([('bbox', {'AP': 31.27763852613224, 'AP50': 50.28872537761709, 'AP75': 35.59114212561016, 'APs': 19.592516332434627, 'APM': 40.660917296157365, 'AP1': 66.81196322710581}])
```

## Very slight difference

Note: some optimization that I used is to save the get\_detection\_data list as a global directory. This made training much faster, as it takes a lot longer to load the images once instead of checking the list every time in the lambda function. I also set the SCORE\_THRESH\_TEST to be 0.75

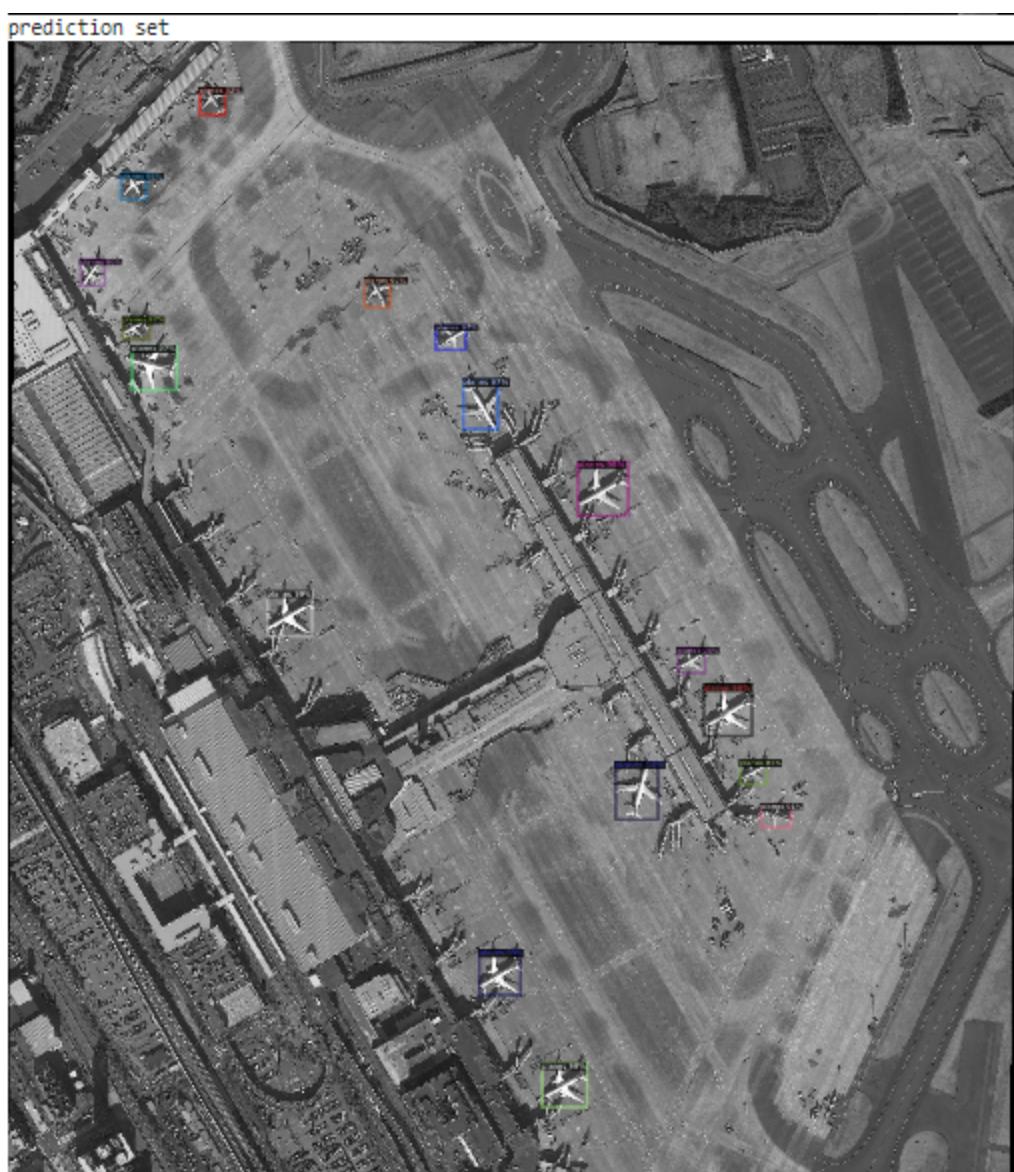
## Example test prediction visualization

### Config:

```
cfg.SOLVERIMS_PER_BATCH = 2 # This is the real "batch size" commonly  
known to deep learning people  
cfg.SOLVER.BASE_LR = 0.0009 # pick a good LR  
cfg.SOLVER.MAX_ITER = 300
```

### Visualization:







Visualization with the following config:

```
cfg.SOLVER.IMS_PER_BATCH = 2 # This is the real "batch size" commonly  
known to deep learning people  
cfg.SOLVER.BASE_LR = 0.000000001 # pick a good LR
```

```
cfg.SOLVER.MAX_ITER = 300
```

Visualization:



`prediction set`



`prediction set`

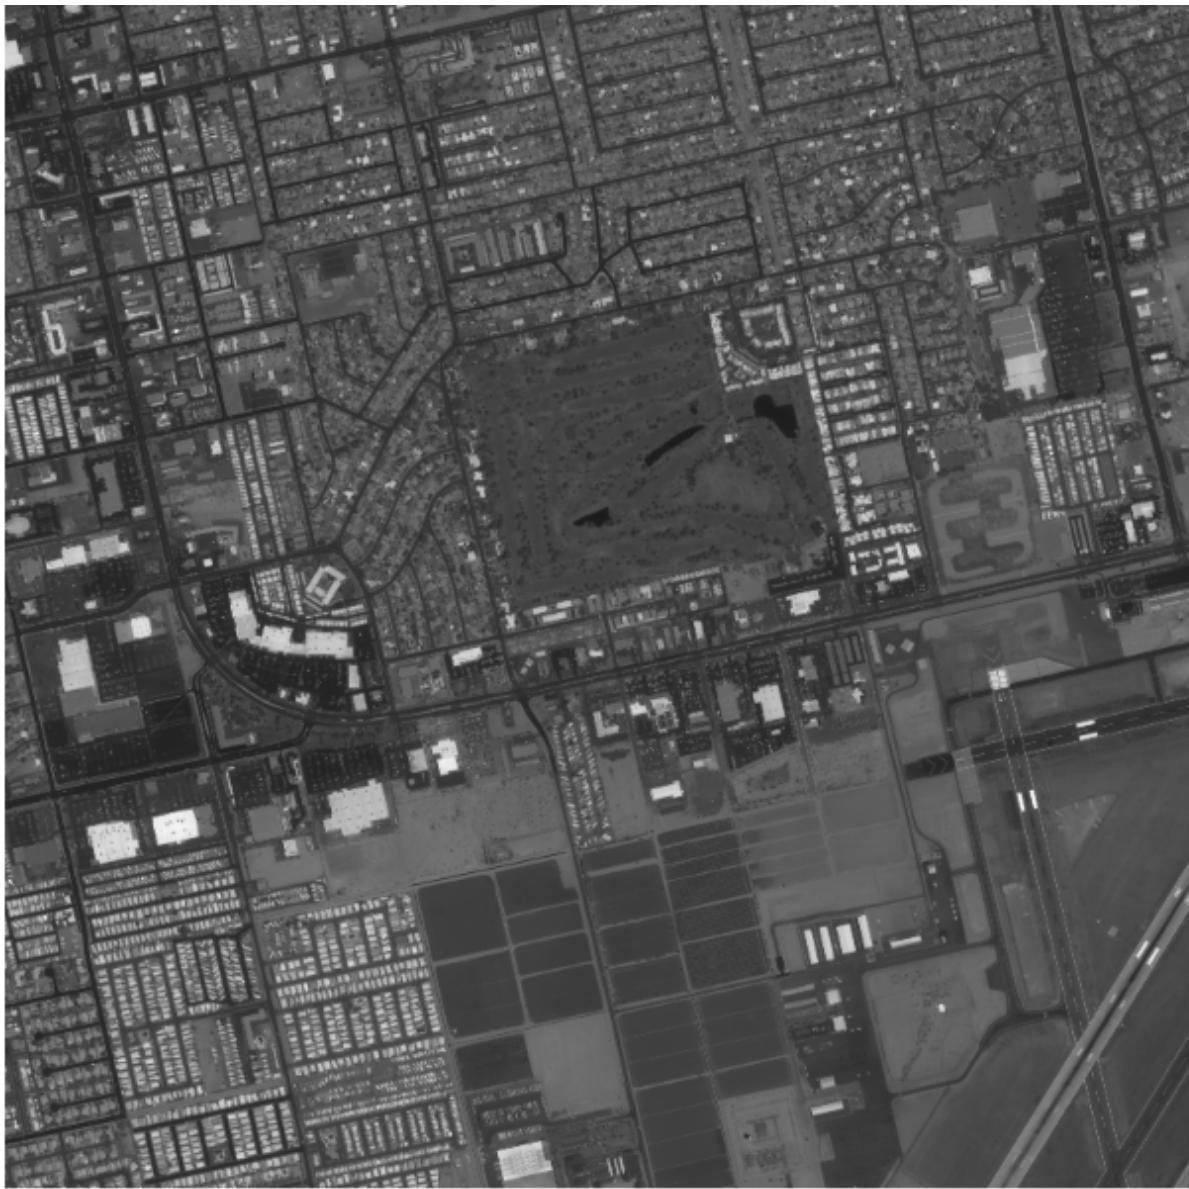


Final Visualization:

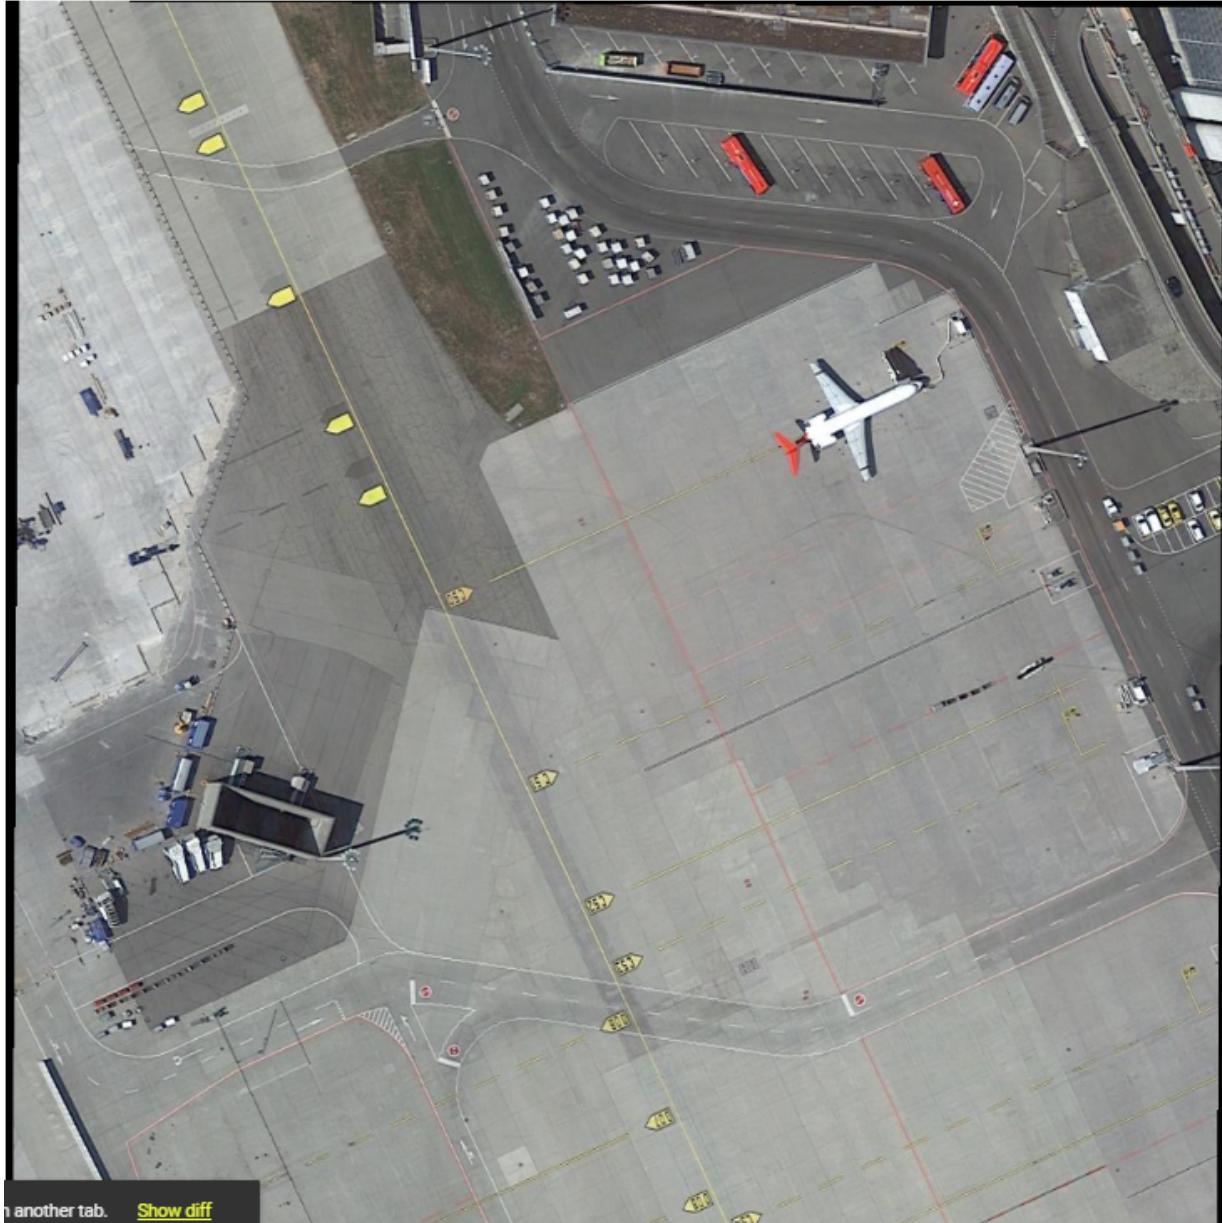
Test set:



Prediction set

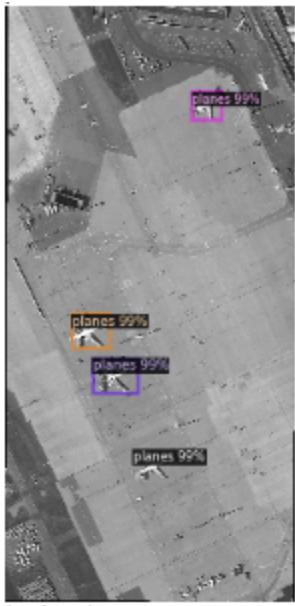


test set

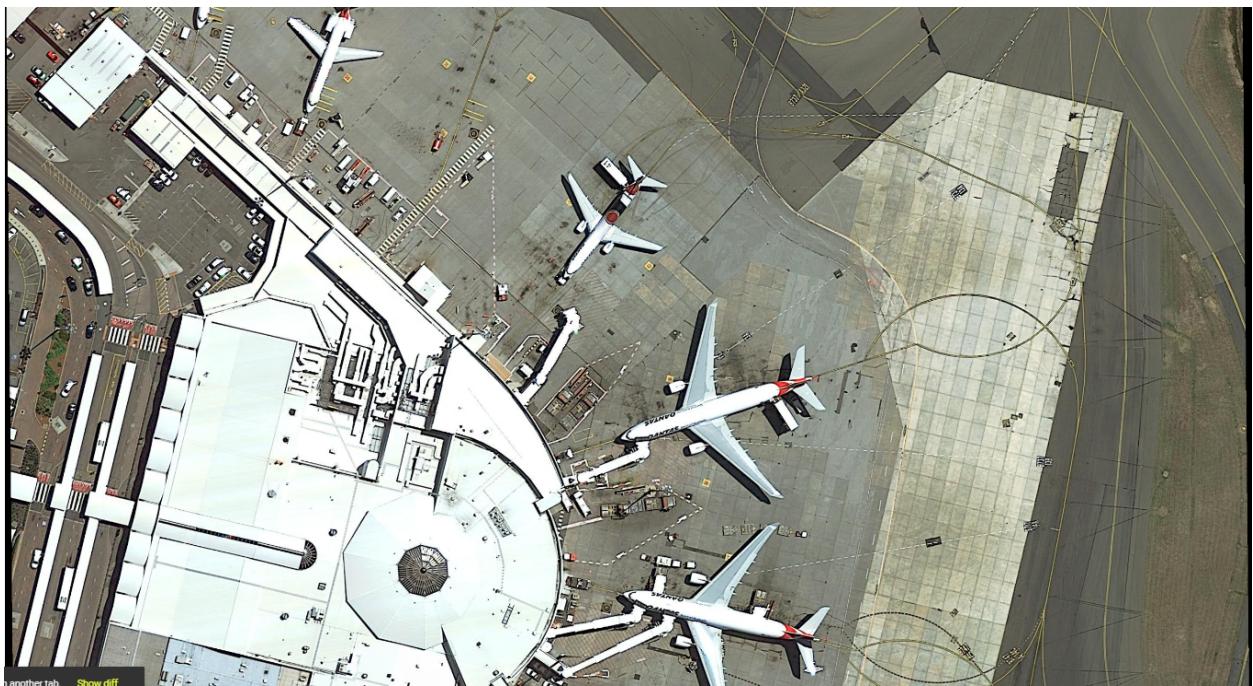


in another tab. [Show diff](#)

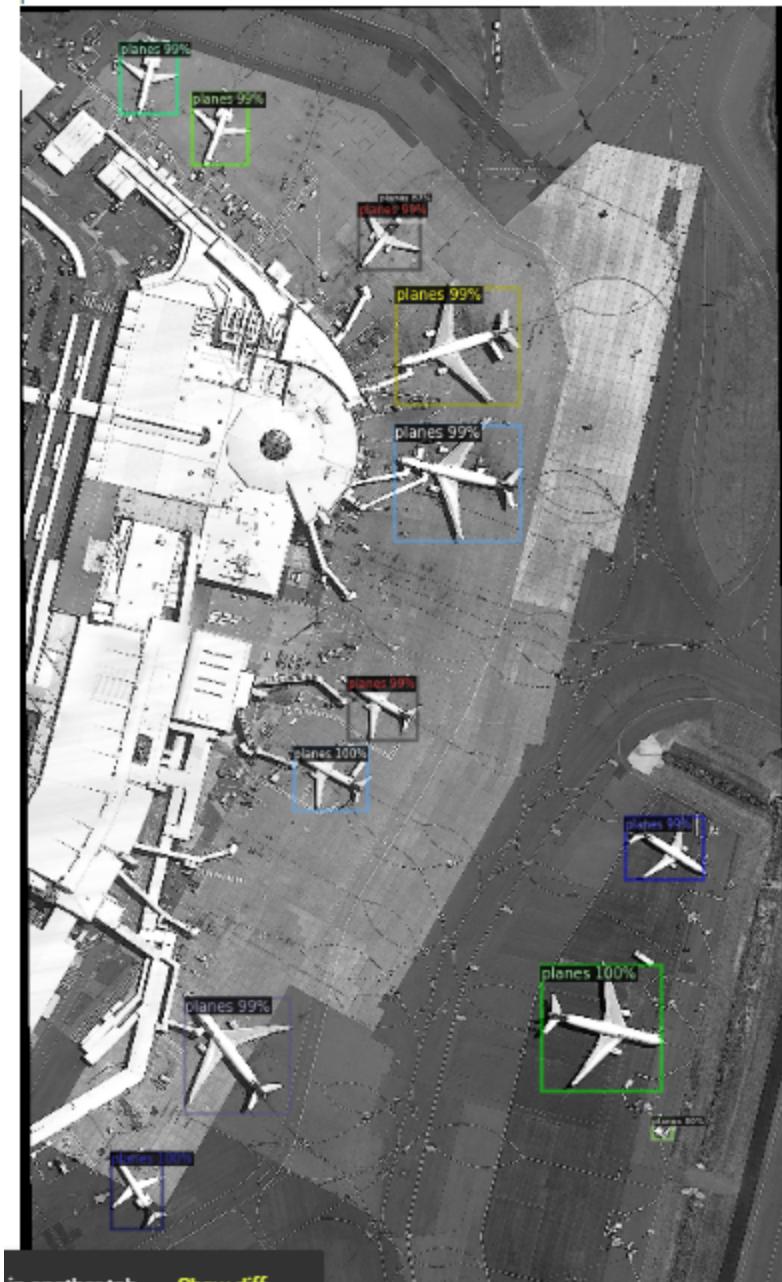
Prediction set



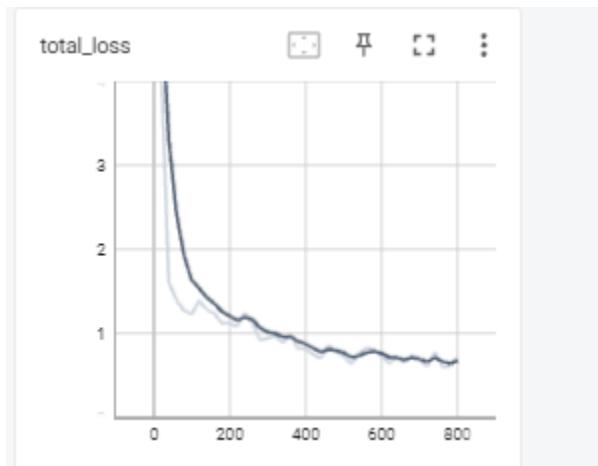
Test set:



Prediction set:



Total loss graph:



## Part 2

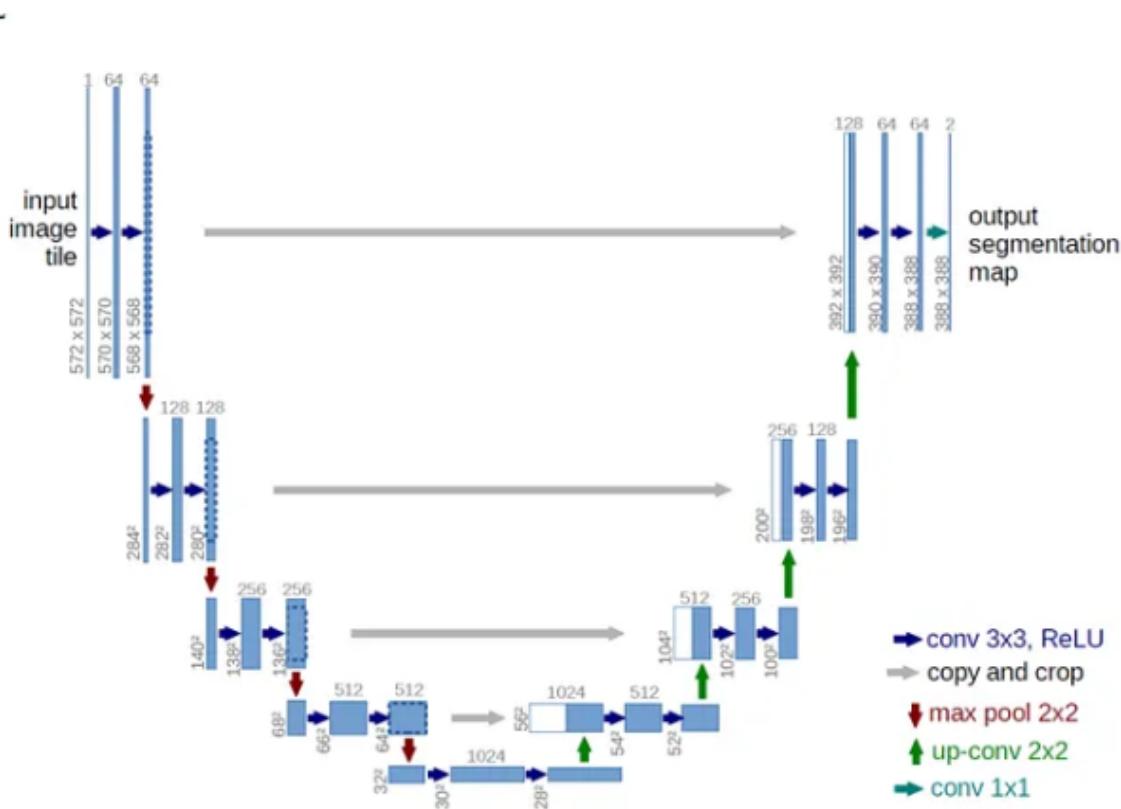
Edits and changes in base code:

Changing the function `get_instance_sample` to load all file names at once instead of by idx individually significantly increased the training time of the Unet model and allowed the function to train up to 50 epochs in 1 hr instead of 1 epoch per hour. This method is memoization. I also pre-loaded all the files as a dict to save on RAM.

Unet model with 1024x1024 as lowest channel visual results:



- note that the prediction is significantly lower ei as shown in the model below:



My modified version of Unet:

Layer (type)	Output Shape	Param #
<hr/>		
Conv2d-1	[ -1, 8, 128, 128]	224
BatchNorm2d-2	[ -1, 8, 128, 128]	16
ReLU-3	[ -1, 8, 128, 128]	0
conv-4	[ -1, 8, 128, 128]	0
Conv2d-5	[ -1, 16, 128, 128]	1,168
BatchNorm2d-6	[ -1, 16, 128, 128]	32
ReLU-7	[ -1, 16, 128, 128]	0
conv-8	[ -1, 16, 128, 128]	0
MaxPool2d-9	[ -1, 16, 64, 64]	0
down-10	[ -1, 16, 64, 64]	0
Conv2d-11	[ -1, 32, 64, 64]	4,640
BatchNorm2d-12	[ -1, 32, 64, 64]	64
ReLU-13	[ -1, 32, 64, 64]	0
conv-14	[ -1, 32, 64, 64]	0
MaxPool2d-15	[ -1, 32, 32, 32]	0
down-16	[ -1, 32, 32, 32]	0
Conv2d-17	[ -1, 64, 32, 32]	18,496
BatchNorm2d-18	[ -1, 64, 32, 32]	128
ReLU-19	[ -1, 64, 32, 32]	0
conv-20	[ -1, 64, 32, 32]	0
MaxPool2d-21	[ -1, 64, 16, 16]	0
down-22	[ -1, 64, 16, 16]	0
Conv2d-23	[ -1, 128, 16, 16]	73,856
BatchNorm2d-24	[ -1, 128, 16, 16]	256
ReLU-25	[ -1, 128, 16, 16]	0
conv-26	[ -1, 128, 16, 16]	0
MaxPool2d-27	[ -1, 128, 8, 8]	0
down-28	[ -1, 128, 8, 8]	0
Conv2d-29	[ -1, 256, 8, 8]	295,168
BatchNorm2d-30	[ -1, 256, 8, 8]	512
ReLU-31	[ -1, 256, 8, 8]	0
conv-32	[ -1, 256, 8, 8]	0
MaxPool2d-33	[ -1, 256, 4, 4]	0
down-34	[ -1, 256, 4, 4]	0
Conv2d-35	[ -1, 512, 4, 4]	1,180,160
BatchNorm2d-36	[ -1, 512, 4, 4]	1,024
ReLU-37	[ -1, 512, 4, 4]	0
conv-38	[ -1, 512, 4, 4]	0
MaxPool2d-39	[ -1, 512, 2, 2]	0
down-40	[ -1, 512, 2, 2]	0
ConvTranspose2d-41	[ -1, 512, 4, 4]	1,049,088
Conv2d-42	[ -1, 256, 4, 4]	1,179,904
BatchNorm2d-43	[ -1, 256, 4, 4]	512
ReLU-44	[ -1, 256, 4, 4]	0
conv-45	[ -1, 256, 4, 4]	0
up-46	[ -1, 256, 4, 4]	0
BatchNorm2d-47	[ -1, 256, 4, 4]	512
ConvTranspose2d-48	[ -1, 256, 8, 8]	262,400
Conv2d-49	[ -1, 128, 8, 8]	295,040
BatchNorm2d-50	[ -1, 128, 8, 8]	256
ReLU-51	[ -1, 128, 8, 8]	0
conv-52	[ -1, 128, 8, 8]	0
up-53	[ -1, 128, 8, 8]	0
BatchNorm2d-54	[ -1, 128, 8, 8]	256
ConvTranspose2d-55	[ -1, 128, 16, 16]	65,664

```

Conv2d-56      [-1, 64, 16, 16]      73,792
BatchNorm2d-57 [-1, 64, 16, 16]      128
    ReLU-58      [-1, 64, 16, 16]      0
    conv-59      [-1, 64, 16, 16]      0
    up-60      [-1, 64, 16, 16]      0
BatchNorm2d-61 [-1, 64, 16, 16]      128
onvTranspose2d-62 [-1, 64, 32, 32]  16,448
    Conv2d-63      [-1, 32, 32, 32]  18,464
BatchNorm2d-64 [-1, 32, 32, 32]      64
    ReLU-65      [-1, 32, 32, 32]      0
    conv-66      [-1, 32, 32, 32]      0
    up-67      [-1, 32, 32, 32]      0
BatchNorm2d-68 [-1, 32, 32, 32]      64
onvTranspose2d-69 [-1, 32, 64, 64]  4,128
    Conv2d-70      [-1, 16, 64, 64]  4,624
BatchNorm2d-71 [-1, 16, 64, 64]      32
    ReLU-72      [-1, 16, 64, 64]      0
    conv-73      [-1, 16, 64, 64]      0
    up-74      [-1, 16, 64, 64]      0
BatchNorm2d-75 [-1, 16, 64, 64]      32
onvTranspose2d-76 [-1, 16, 128, 128] 1,040
    Conv2d-77      [-1, 8, 128, 128] 1,160
BatchNorm2d-78 [-1, 8, 128, 128]      16
    ReLU-79      [-1, 8, 128, 128]      0
    conv-80      [-1, 8, 128, 128]      0
    up-81      [-1, 8, 128, 128]      0
BatchNorm2d-82 [-1, 8, 128, 128]      16
    Conv2d-83      [-1, 3, 128, 128]  219
BatchNorm2d-84 [-1, 3, 128, 128]      6
    ReLU-85      [-1, 3, 128, 128]      0
    conv-86      [-1, 3, 128, 128]      0
    Conv2d-87      [-1, 1, 128, 128]  28
    conv-88      [-1, 1, 128, 128]      0
=====
al params: 4,549,765
inable params: 4,549,765
-trainable params: 0
-----
ut size (MB): 0.19
ward/backward pass size (MB): 39.22
ams size (MB): 17.36
imated Total Size (MB): 56.76
-----
```

Did not include unsample cause it slowed down the model too much. Also did not go up to 1024x1024 since it causes too much loss of information, as shown in the section above. I also normalized the output of the copy and crop section to eliminate more noise.

Current config:

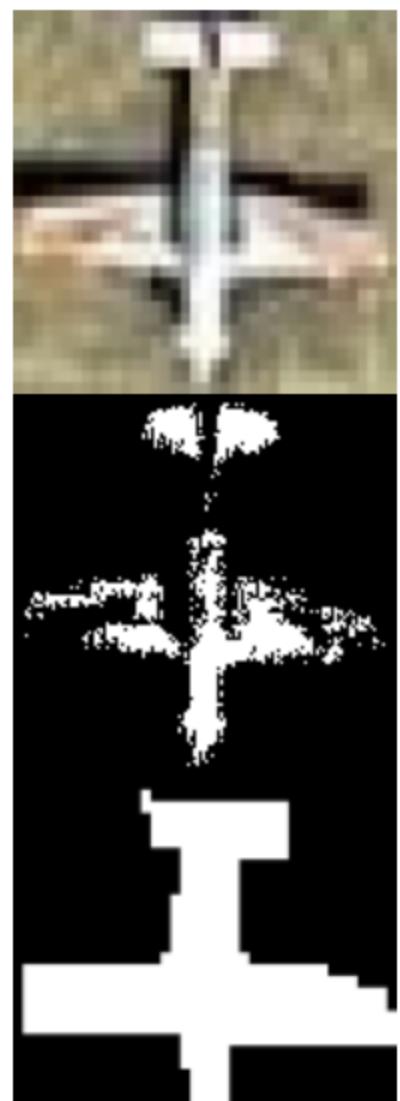
```

num_epochs = 60
batch_size = 8
learning_rate = 0.0001
weight_decay = 1e-5

```

Note: in visualization the order of the photos goes from original image then prediction than true mask.

Resulting visualization:







---

Mean IOU:

#images: 7980, Mean IoU: 0.7003960934377673

Change learning rate to 0.1

Final results in mean IOU:

---

⌚ 100%  998/998 [00:11<00:00, 94.22it/s]

#images: 7980, Mean IoU: 0.9068593285289236

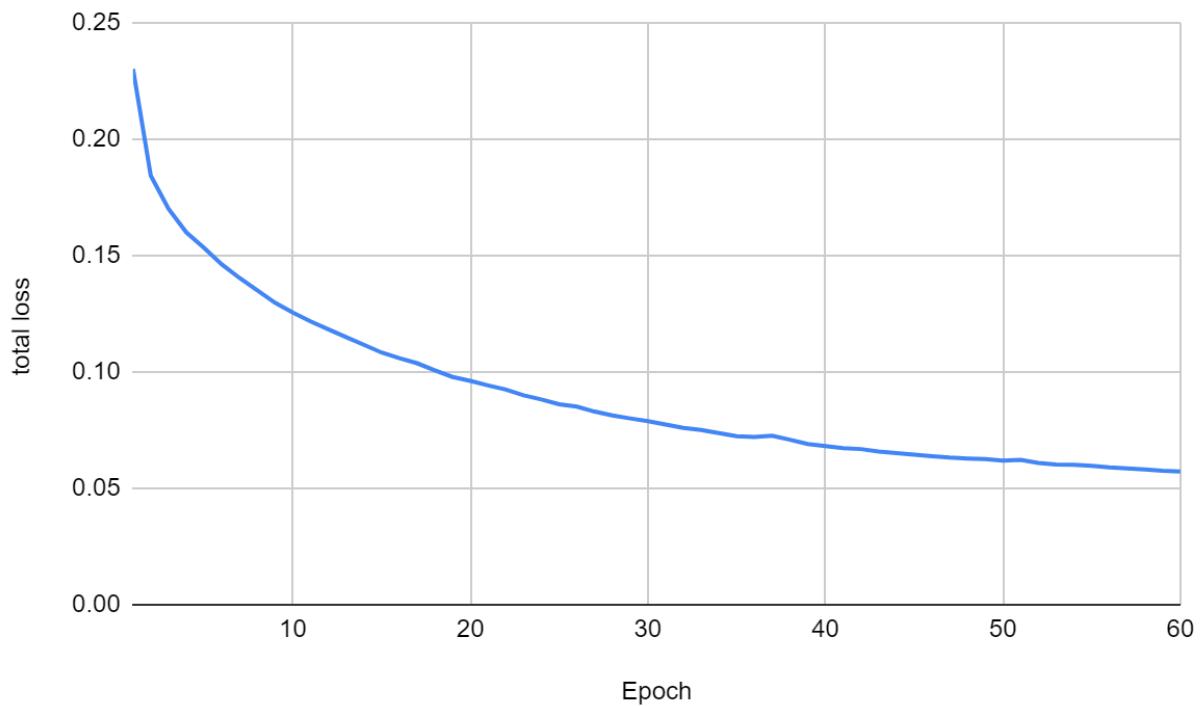
Final Visualization:





Loss function: `BCEWithLogitsLoss()`

Loss graph:



### Part 3:

Group name on kaggle: hrwei\_sfу

Group members: Justin Lin, Hao Ran Wei

Current best score on kaggle: 0.63761

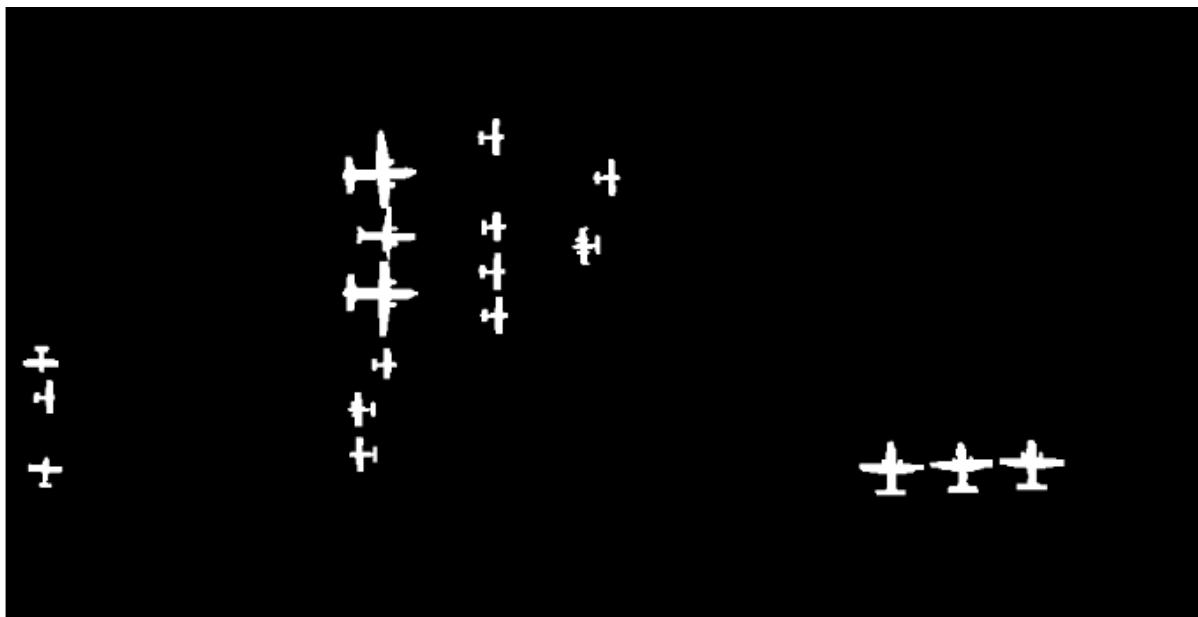
Note: due to the size of the image, the screen shots may not cover the entire image.

Visualization:

original:



True mask



prediction:



original



True mask



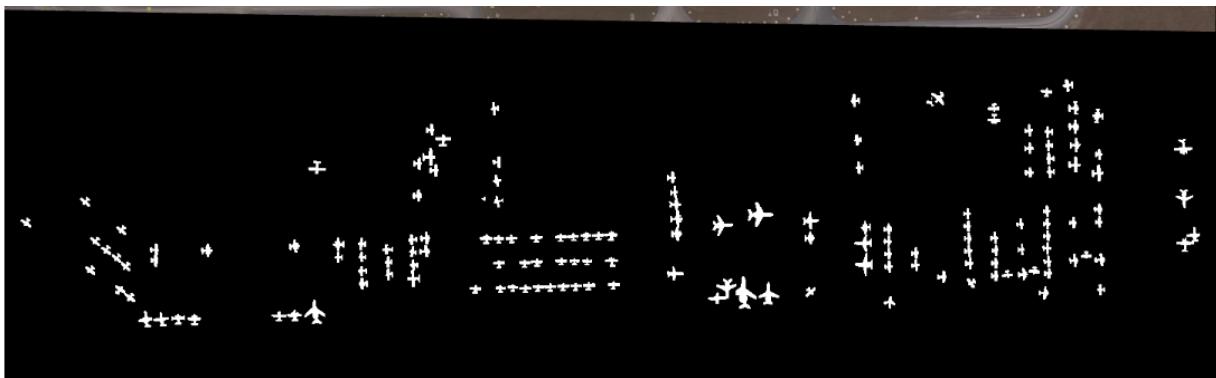
Prediction mask



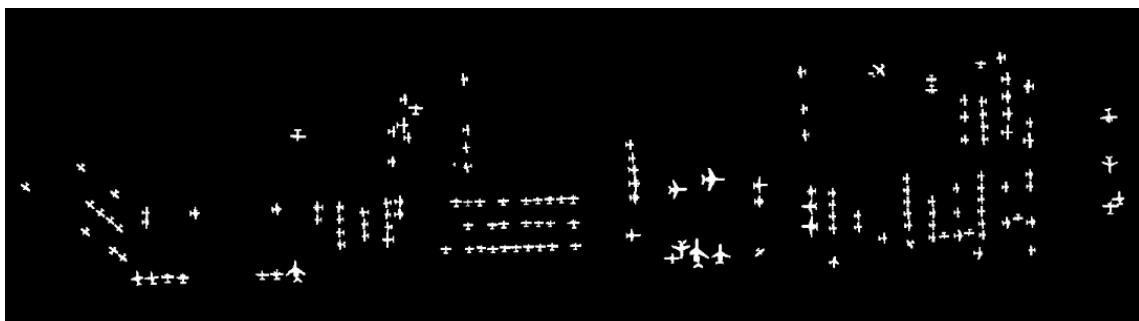
original:



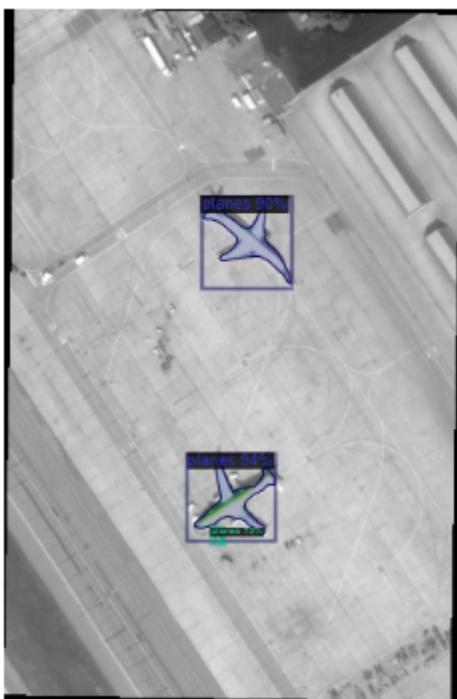
True mask

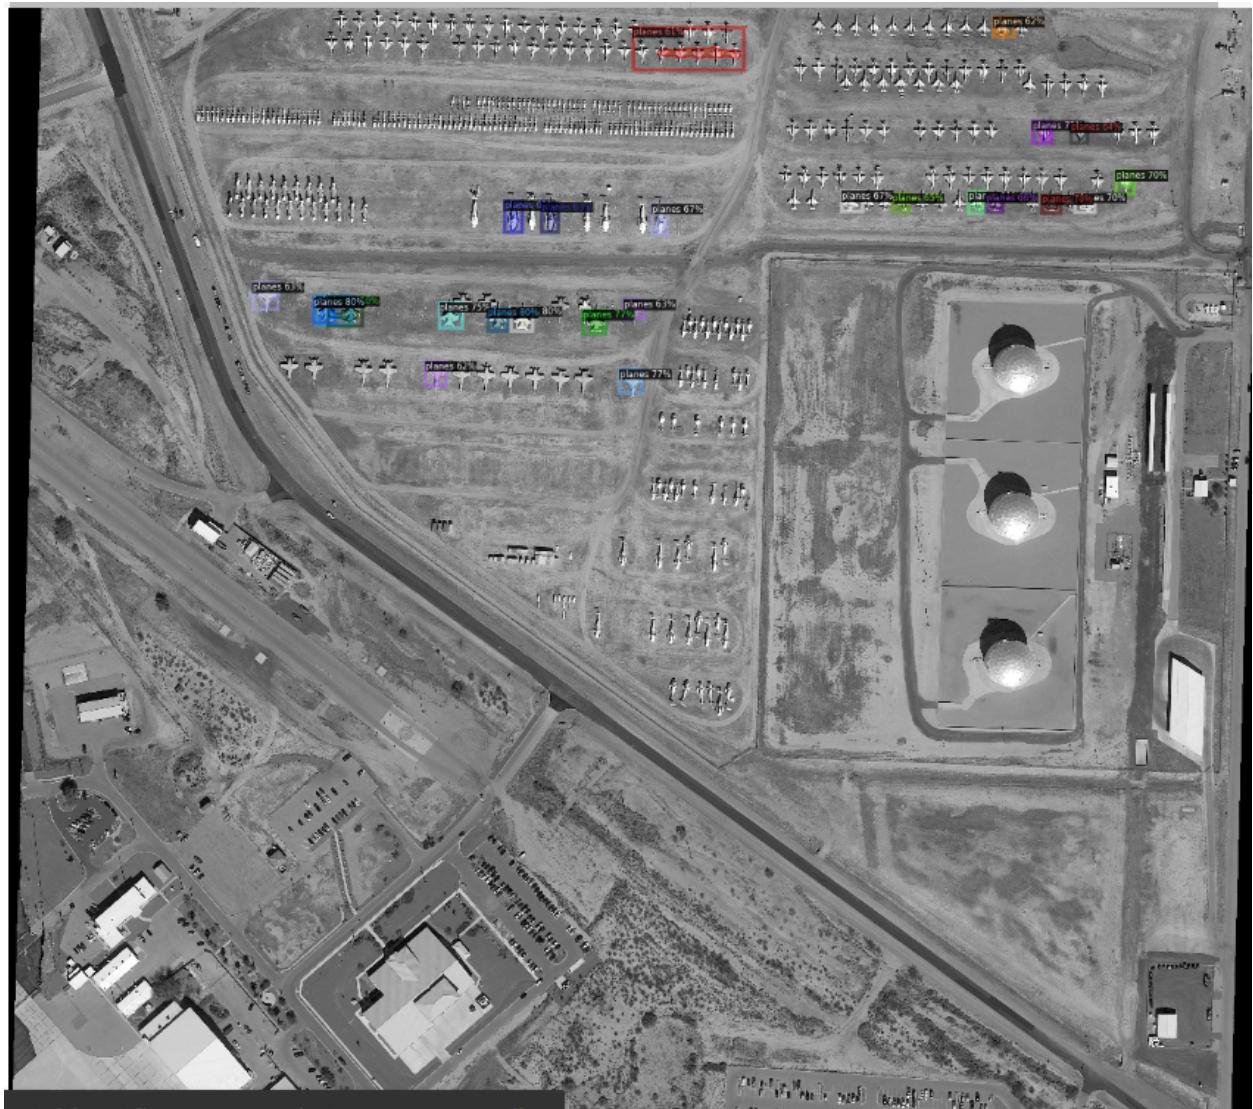


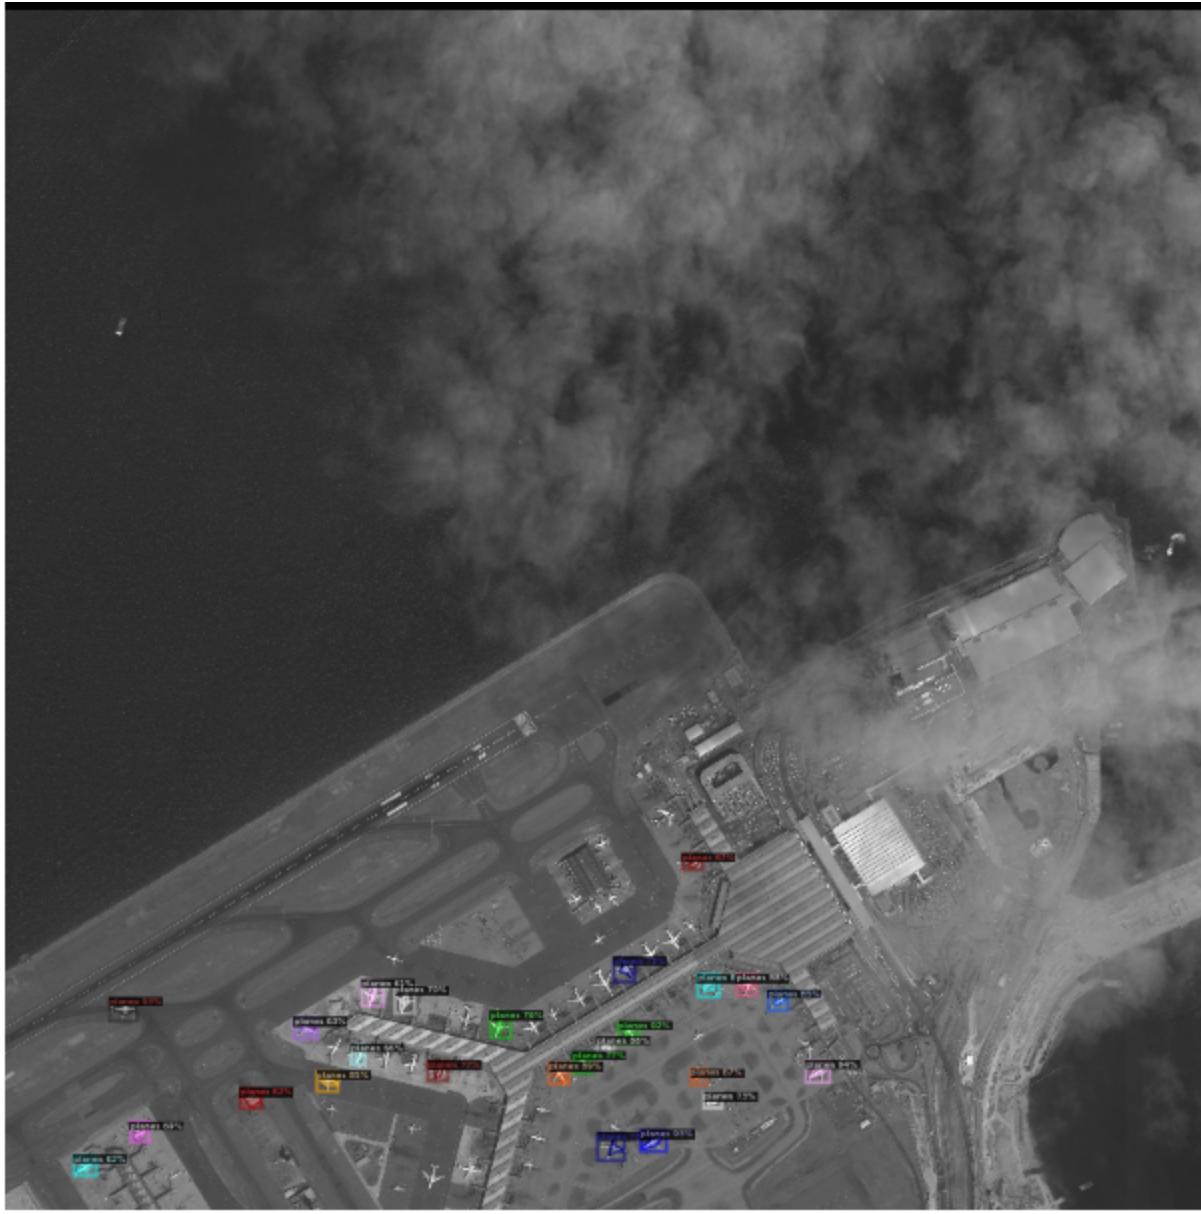
Prediction mask:



Part 4:  
visualization:







This model is way worse than the one used in part 1. The Ap score for this model is 0.119 where as in part 1 it's around 0.31

```
[02/24 02:20:11 d2.evaluation.fast_eval_api]: Evaluate annotation type *segm*
[02/24 02:20:12 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 1.52 seconds.
[02/24 02:20:12 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[02/24 02:20:12 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.03 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.000
[02/24 02:20:12 d2.evaluation.coco_evaluation]: Evaluation results for segm:
| AP      | AP50     | AP75     | APs     | APM     | APl     |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 0.000  | 0.001  | 0.000  | 0.000  | 0.000  | 0.000  |
OrderedDict([('bbox', {'AP': 0.11915358397089167, 'AP50': 0.532702566638666, 'AP75': 0.011977004152028105, 'APs': 0.04021587619959273, 'APM': 0.276})
```